

Lab 3: Radial Basis Function (RBF) Networks

This notebook provides a basic, guided walkthrough of Radial Basis Function networks. You can run all the code and study how it works, then experiment with the parameters.

Learning Objectives:

1. Understand the architecture of RBF networks
 2. See how RBF differs from Multi-Layer Perceptrons (MLPs)
 3. Implement an RBF network from scratch
 4. Apply RBF to the Iris classification problem
 5. Compare RBF with MLP performance
 6. Use sklearn's kernel methods as a library alternative
-

1. Introduction: What is an RBF Network?


An **RBF network** is a three-layer neural network used for:

- Classification
- Function approximation
- Time series prediction
- System control

Key Differences from MLP

Aspect	MLP	RBF
Hidden layer activation	Sigmoid (global)	Gaussian (local)
Hidden layer weights	Learned via backpropagation	None - stores prototypes
Output layer activation	Sigmoid	None (linear)
Training	Iterative gradient descent	Two-stage (prototypes, then weights)
Response	All neurons respond to all inputs	Only nearby neurons respond strongly

RBF Architecture



Each RBF neuron stores a prototype (p_i) and computes similarity between the input and its prototype using a Gaussian function.

Key insight: RBF neurons respond strongly only when the input is *close* to their prototype. This is fundamentally different from MLP neurons, which respond to weighted sums across the entire input space.

In []: # Import required libraries

```
# Avoid KMeans memory leak issue on Windows with Intel MKL
import os
os.environ['OMP_NUM_THREADS'] = '1'
import warnings
warnings.filterwarnings('ignore', message='KMeans is known to have a memory leak')

import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.cluster import KMeans
from sklearn.metrics import accuracy_score, confusion_matrix

np.random.seed(42)
print("Libraries loaded successfully!")
```

Libraries loaded successfully!

2. The Gaussian (RBF) Activation Function

The core of an RBF network is the **Gaussian activation function**:

$$\varphi(x) = e^{-\beta \cdot \|x - p\|^2}$$

Where:

- x is the input vector
- p is the prototype (centre) stored by the neuron
- $\|x - p\|$ is the Euclidean distance between input and prototype
- β controls the width of the Gaussian curve

The β Parameter

- **Large β** → Narrow curve → Neuron responds only to inputs very close to prototype
- **Small β** → Wide curve → Neuron responds to inputs further from prototype

β can be computed from the standard deviation (σ) of the data:

$$\beta = \frac{1}{2\sigma^2}$$

In []: # Visualise the Gaussian function with different β values

```
def gaussian(x, centre, beta):
    """
    Compute Gaussian activation.

    Parameters:
        x: Input value(s)
        centre: The prototype/centre of the Gaussian
        beta: Width parameter (larger = narrower curve)
```

```

>Returns:
Activation value(s) in range [0, 1]
"""
distance_squared = (x - centre) ** 2
## TODO: use the numpy's "exp" function to return the gaussian activation
##         function value (see equation in the markdown cell above)

# Create x values
## TODO: use the numpy's "Linspace" function to create 200 evenly
##       spaced values for x over the interval [-4,4].
centre = 0 # Prototype at origin


# Plot Gaussians with different β values
plt.figure(figsize=(10, 6))

for beta, color, label in [(0.5, 'blue', 'β = 0.5 (wide)'),
                           (1.0, 'orange', 'β = 1.0 (medium)'),
                           (2.0, 'red', 'β = 2.0 (narrow)')]:
    y = gaussian(x, centre, beta)
    plt.plot(x, y, color=color, linewidth=2, label=label)

plt.axhline(y=0.5, color='grey', linestyle='--', alpha=0.5)
plt.axvline(x=0, color='grey', linestyle='--', alpha=0.5)
plt.xlabel('Distance from prototype', fontsize=12)
plt.ylabel('Activation φ(x)', fontsize=12)
plt.title('Gaussian Activation Function with Different β Values', fontsize=14)
plt.legend(fontsize=11)
plt.grid(True, alpha=0.3)
plt.ylim(-0.05, 1.05)
plt.tight_layout()
plt.savefig('gaussian_activation.png', dpi=150)
plt.show()

print("\nKey observation:")
print("- When input equals prototype (distance = 0), activation = 1.0")
print("- As input moves away from prototype, activation falls towards 0")
print("- β controls how quickly the activation falls off")

```



Key observation:

- When input equals prototype ($distance = 0$), activation = 1.0
 - As input moves away from prototype, activation falls towards 0
 - β controls how quickly the activation falls off
-

3. RBF Network Implementation (From Scratch)

Now we'll implement a complete RBF network. The training process has two stages:

1. **Select prototypes** - Choose which training samples (or cluster centres) to use as RBF neuron centres
2. **Learn output weights** - Find the weights that map RBF activations to correct outputs

Prototype Selection Methods

From the lecture slides, three options:

1. One RBF neuron for each training sample (expensive)
2. Randomly select k prototypes from the dataset
3. Use k-means clustering and take cluster centroids as prototypes (recommended)

We'll use **k-means clustering** as it provides well-distributed prototypes.

```
In [ ]: class RBFNetwork:  
    """  
        Radial Basis Function Network for classification.  
  
        Architecture:  
            - Input layer: passes input directly to hidden layer  
            - Hidden layer: RBF neurons, each storing a prototype  
            - Output layer: weighted sum (no activation function)  
  
        Training:  
            1. Select prototypes using k-means clustering  
            2. Compute  $\beta$  from cluster spreads  
            3. Learn output weights using pseudoinverse  
    """  
  
    def __init__(self, n_prototypes=10):  
        """  
            Initialise RBF network.  
  
            Parameters:  
                n_prototypes: Number of RBF neurons (prototypes) in hidden layer  
        """  
        self.n_prototypes = n_prototypes  
        self.prototypes = None # Will store prototype vectors  
        self.betas = None # Will store  $\beta$  for each prototype  
        self.weights = None # Will store output layer weights  
  
    def _euclidean_distance(self, x, prototype):  
        """  
            Compute Euclidean distance between input and prototype.  
  
             $||x - p|| = \sqrt{\sum(x_i - p_i)^2}$   
        """  
        ## TODO: use numpy's "sqrt" and "sum" (making axis=-1) functions  
        ##      to compute and return the Euclidean distance
```

```

def _gaussian_activation(self, distance, beta):
    """
    Compute Gaussian activation:  $\phi = \exp(-\beta * d^2)$ 

    Parameters:
        distance: Euclidean distance from prototype
        beta: Width parameter

    Returns:
        Activation value in [0, 1]
    """
    return np.exp(-beta * distance ** 2)

def _compute_hidden_activations(self, X):
    """
    Compute activations of all RBF neurons for all inputs.

    Parameters:
        X: Input data, shape (n_samples, n_features)

    Returns:
        Activation matrix, shape (n_samples, n_prototypes)
    """
    ## TODO: given that 'X' is a 2D NumPy array, read its number of
    ##       rows (i.e. samples) into a new variable named "n_samples".

    activations = np.zeros((n_samples, self.n_prototypes))

    # For each prototype, compute activation for all samples
    for j, (prototype, beta) in enumerate(zip(self.prototypes, self.betas)):
        distances = self._euclidean_distance(X, prototype)
        activations[:, j] = self._gaussian_activation(distances, beta)

    return activations

def fit(self, X, y):
    """
    Train the RBF network.

    Stage 1: Select prototypes using k-means clustering
    Stage 2: Compute  $\beta$  values from cluster spreads
    Stage 3: Learn output weights using least squares

    Parameters:
        X: Training data, shape (n_samples, n_features)
        y: Target labels (one-hot encoded), shape (n_samples, n_classes)
    """
    print(f"Training RBF Network with {self.n_prototypes} prototypes...")

    # =====
    # STAGE 1: Select prototypes using k-means clustering
    # =====
    # K-means finds cluster centres that represent the data distribution
    # These centres become our RBF neuron prototypes

    kmeans = KMeans(n_clusters=self.n_prototypes, random_state=42, n_init=10)
    kmeans.fit(X)
    self.prototypes = kmeans.cluster_centers_

    print(f"  Stage 1: Selected {self.n_prototypes} prototypes via k-means")

    # =====
    # STAGE 2: Compute  $\beta$  for each prototype
    # =====
    #  $\beta = 1 / (2\sigma^2)$  where  $\sigma$  is the spread of the cluster
    # We use the average distance from cluster centre to its members

```

```

    self.betas = np.zeros(self.n_prototypes)

    for j in range(self.n_prototypes):
        # Find all points belonging to this cluster
        cluster_points = X[kmeans.labels_ == j]

        if len(cluster_points) > 1:
            # Compute average distance to cluster centre
            distances = self._euclidean_distance(cluster_points, self.prototypes[j])
            ## TODO: use NumPy's "mean" function to calculate "sigma" as the mean
            ##       of the Euclidean "distances"

            # Avoid division by zero
            if sigma > 0:
                self.betas[j] = 1 / (2 * sigma ** 2)
            else:
                self.betas[j] = 1.0 # Default if all points identical
        else:
            self.betas[j] = 1.0 # Default for single-point clusters

    print(f" Stage 2: Computed β values (mean={self.betas.mean():.3f})")

# =====
# STAGE 3: Learn output weights using least squares
# =====
# We need to find weights W such that: H @ W ≈ Y
# Where H is the hidden layer activation matrix
# Solution: W = (H^T H)^(-1) H^T Y (pseudoinverse)

# Compute hidden layer activations for all training samples
H = self._compute_hidden_activations(X)

# Add bias column (column of 1s) to handle offset
H_bias = np.column_stack([H, np.ones(H.shape[0])])

# Compute weights using pseudoinverse (more stable than direct inverse)
# W = pinv(H) @ Y
self.weights = np.linalg.pinv(H_bias) @ y

print(f" Stage 3: Learned output weights, shape={self.weights.shape}")
print("Training complete!")

return self

def predict_proba(self, X):
    """
    Predict class probabilities for input samples.

    Parameters:
        X: Input data, shape (n_samples, n_features)

    Returns:
        Class scores, shape (n_samples, n_classes)
    """

    # Compute hidden layer activations
    H = self._compute_hidden_activations(X)

    # Add bias column
    H_bias = np.column_stack([H, np.ones(H.shape[0])])

    # Compute output: weighted sum (no activation function)
    output = H_bias @ self.weights

    return output

```

```

def predict(self, X):
    """
    Predict class labels for input samples.

    Parameters:
        X: Input data, shape (n_samples, n_features)

    Returns:
        Predicted class indices, shape (n_samples,)
    """
    scores = self.predict_proba(X)
    ## TODO: use NumPy's "argmax" function to identify and return the
    ##       indices of the highest values for each row of "scores".

```

```

print("RBFNetwork class defined successfully!")
print("\nKey components:")
print(" - _euclidean_distance(): Computes ||x - p||")
print(" - _gaussian_activation(): Computes exp(-β * d²)")
print(" - fit(): Three-stage training (prototypes, betas, weights)")
print(" - predict(): Classification using trained network")

```

RBFNetwork class defined successfully!

Key components:

- `_euclidean_distance()`: Computes $||x - p||$
 - `_gaussian_activation()`: Computes $\exp(-\beta \cdot d^2)$
 - `fit()`: Three-stage training (prototypes, betas, weights)
 - `predict()`: Classification using trained network
-

4. Apply to Iris Dataset

Now let's apply our RBF network to the Iris dataset - the same dataset we used in Lab 2 with the MLP.

In []: # Load and prepare the Iris dataset

```

# Load data
iris = load_iris()
X = iris.data
y = iris.target

print("Iris Dataset")
print("=" * 40)
print(f"Features: {iris.feature_names}")
print(f"Classes: {iris.target_names}")
print(f"Samples: {X.shape[0]}, Features: {X.shape[1]}")

# Normalise features to [0, 1] range
scaler = MinMaxScaler()
X_normalised = scaler.fit_transform(X)

# One-hot encode Labels
n_classes = len(np.unique(y))
y_onehot = np.zeros((len(y), n_classes))
for i, label in enumerate(y):
    y_onehot[i, label] = 1

print(f"\nNormalised feature range: [{X_normalised.min():.2f}, {X_normalised.max():.2f}]")
print(f"One-hot encoded labels shape: {y_onehot.shape}")

```

```

# Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(
    X_normalised, y_onehot, test_size=0.2, random_state=42, stratify=y
)
y_test_labels = np.argmax(y_test, axis=1)

print(f"\nTraining samples: {X_train.shape[0]}")
print(f"Test samples: {X_test.shape[0]}")

```

Iris Dataset

```

=====
Features: ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
Classes: ['setosa' 'versicolor' 'virginica']
Samples: 150, Features: 4

```

```

Normalised feature range: [0.00, 1.00]
One-hot encoded labels shape: (150, 3)

```

Training samples: 120

Test samples: 30

In []: # Train the RBF network

```

# For Iris with 3 classes, we'll use 10 prototypes per class = 30 total
# This follows the slide example: "e.g. 10 samples/class → 3 * 10 = 30 RBF neurons"
n_prototypes = 30

rbf = RBFNetwork(n_prototypes=n_prototypes)
rbf.fit(X_train, y_train);

```

Training RBF Network with 30 prototypes...

Stage 1: Selected 30 prototypes via k-means

Stage 2: Computed β values (mean=90.610)

Stage 3: Learned output weights, shape=(31, 3)

Training complete!

In []: # Evaluate the RBF network

```

# Predictions
y_train_pred = rbf.predict(X_train)
y_test_pred = rbf.predict(X_test)

# Accuracy
train_accuracy = accuracy_score(np.argmax(y_train, axis=1), y_train_pred)
test_accuracy = accuracy_score(y_test_labels, y_test_pred)

print("RBF Network Results")
print("=" * 40)
print(f"Training Accuracy: {train_accuracy * 100:.2f}%")
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")

# Confusion matrix
cm = confusion_matrix(y_test_labels, y_test_pred)
print(f"\nConfusion Matrix:")
print(cm)

# Visualise confusion matrix
fig, ax = plt.subplots(figsize=(8, 6))
im = ax.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
ax.figure.colorbar(im, ax=ax)
ax.set(xticks=np.arange(cm.shape[1]),
       yticks=np.arange(cm.shape[0]),
       xticklabels=iris.target_names,
       yticklabels=iris.target_names,
       title='RBF Network - Confusion Matrix',
       ylabel='Actual',
       xlabel='Predicted')

```

```

        xlabel='Predicted')

# Add text annotations
thresh = cm.max() / 2.
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, format(cm[i, j], 'd'),
                ha="center", va="center",
                color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.savefig('rbf_confusion_matrix.png', dpi=150)
plt.show()

```

RBF Network Results


=====

Training Accuracy: 96.67%

Test Accuracy: 93.33%

Confusion Matrix:

```
[[10  0  0]
 [ 0 10  0]
 [ 0  2  8]]
```



Visualise RBF Neuron Activations

Let's see how individual RBF neurons respond to different input samples. This illustrates the "local" nature of RBF - each neuron responds strongly only to inputs similar to its prototype.

In []: # Visualise RBF activations for sample inputs

```

# Compute activations for test set
activations = rbf._compute_hidden_activations(X_test)

# Select 3 samples (one from each class)
sample_indices = []
for c in range(3):
    idx = np.where(y_test_labels == c)[0][0]
    sample_indices.append(idx)


# Plot activation patterns
fig, axes = plt.subplots(1, 3, figsize=(15, 4))

for i, (idx, ax) in enumerate(zip(sample_indices, axes)):
    ax.bar(range(rbf.n_prototypes), activations[idx], color='steelblue', alpha=0.7)
    ax.set_xlabel('RBF Neuron (Prototype) Index')
    ax.set_ylabel('Activation')
    ax.set_title(f'{iris.target_names[y_test_labels[idx]]} Sample')
    ax.set_ylim(0, 1.05)
    ax.axhline(y=0.5, color='red', linestyle='--', alpha=0.5, label='0.5 threshold')
    ax.legend()

plt.tight_layout()
plt.savefig('rbf_activations.png', dpi=150)
plt.show()

print("\nObservation:")
print("Each sample activates only a few RBF neurons strongly (those with similar prototypes).")
print("This is the 'local response' property of RBF networks.")

```



Observation:

Each sample activates only a few RBF neurons strongly (those with similar prototypes). This is the 'local response' property of RBF networks.

5. Compare with MLP

Let's compare our RBF network with an MLP on the same dataset. We'll use our MLP implementation from Lab 2.

In []: # Simple MLP implementation (from Lab 2)

```

class MLP:
    """Multi-Layer Perceptron with one hidden layer."""

    def __init__(self, n_inputs, n_hidden, n_outputs, learning_rate=0.5):
        self.learning_rate = learning_rate

        # Initialise weights randomly in [-0.5, 0.5]
        self.weights_ih = np.random.uniform(-0.5, 0.5, (n_inputs, n_hidden))
        self.bias_h = np.random.uniform(-0.5, 0.5, n_hidden)
        self.weights_ho = np.random.uniform(-0.5, 0.5, (n_hidden, n_outputs))
        self.bias_o = np.random.uniform(-0.5, 0.5, n_outputs)

```

```

def sigmoid(self, x):
    return 1 / (1 + np.exp(-np.clip(x, -500, 500)))

def forward(self, X):
    self.z_h = np.dot(X, self.weights_ih) + self.bias_h
    self.a_h = self.sigmoid(self.z_h)
    self.z_o = np.dot(self.a_h, self.weights_ho) + self.bias_o
    self.a_o = self.sigmoid(self.z_o)
    return self.a_o

def backward(self, X, y):
    n_samples = X.shape[0]
    output_error = y - self.a_o
    output_delta = output_error * self.a_o * (1 - self.a_o)
    hidden_error = np.dot(output_delta, self.weights_ho.T)
    hidden_delta = hidden_error * self.a_h * (1 - self.a_h)

    self.weights_ho += self.learning_rate * np.dot(self.a_h.T, output_delta) / n_samples
    self.bias_o += self.learning_rate * np.mean(output_delta, axis=0)
    self.weights_ih += self.learning_rate * np.dot(X.T, hidden_delta) / n_samples
    self.bias_h += self.learning_rate * np.mean(hidden_delta, axis=0)

def train(self, X, y, epochs, verbose=False):
    losses = []
    for epoch in range(epochs):
        output = self.forward(X)
        loss = np.mean((y - output) ** 2)
        losses.append(loss)
        self.backward(X, y)
        if verbose and epoch % 200 == 0:
            print(f"Epoch {epoch}: MSE = {loss:.6f}")
    return losses

def predict(self, X):
    output = self.forward(X)
    return np.argmax(output, axis=1)

# Train MLP with same training data
print("Training MLP for comparison...")
print("=" * 40)

np.random.seed(42)
mlp = MLP(n_inputs=4, n_hidden=8, n_outputs=3, learning_rate=0.5)
mlp_losses = mlp.train(X_train, y_train, epochs=1000, verbose=True)

# Evaluate MLP
mlp_train_pred = mlp.predict(X_train)
mlp_test_pred = mlp.predict(X_test)

mlp_train_accuracy = accuracy_score(np.argmax(y_train, axis=1), mlp_train_pred)
mlp_test_accuracy = accuracy_score(y_test_labels, mlp_test_pred)

print(f"\nMLP Training Accuracy: {mlp_train_accuracy * 100:.2f}%")
print(f"MLP Test Accuracy: {mlp_test_accuracy * 100:.2f}%")

```

Training MLP for comparison...

```
=====
Epoch 0: MSE = 0.245520
Epoch 200: MSE = 0.192939
Epoch 400: MSE = 0.134717
Epoch 600: MSE = 0.113445
Epoch 800: MSE = 0.104507
```

MLP Training Accuracy: 90.00%
MLP Test Accuracy: 80.00%

```
In [ ]: # Compare RBF vs MLP
```

```
print("\n" + "=" * 50)
print("COMPARISON: RBF vs MLP on Iris Dataset")
print("=" * 50)

print(f"\n{'Metric':<25} {'RBF':>12} {'MLP':>12}")
print("-" * 50)
print(f"{'Training Accuracy':<25} {train_accuracy*100:>11.2f}% {mlp_train_accuracy*100:>11.2f}%")
print(f"{'Test Accuracy':<25} {test_accuracy*100:>11.2f}% {mlp_test_accuracy*100:>11.2f}%")
print(f"{'Hidden Neurons':<25} {rbf.n_prototypes:>12} {8:>12}")
print(f"{'Training Method':<25} {'Pseudoinverse':>12} {'Backprop':>12}")
print(f"{'Activation (Hidden)':<25} {'Gaussian':>12} {'Sigmoid':>12}")
print(f"{'Activation (Output)':<25} {'Linear':>12} {'Sigmoid':>12}")

# Plot comparison
fig, ax = plt.subplots(figsize=(8, 5))

x_pos = np.arange(2)
width = 0.35

bars1 = ax.bar(x_pos - width/2, [train_accuracy*100, mlp_train_accuracy*100],
                width, label='Training', color='steelblue')
bars2 = ax.bar(x_pos + width/2, [test_accuracy*100, mlp_test_accuracy*100],
                width, label='Test', color='coral')

ax.set_ylabel('Accuracy (%)')
ax.set_title('RBF vs MLP Performance on Iris Dataset')
ax.set_xticks(x_pos)
ax.set_xticklabels(['RBF Network', 'MLP'])
ax.legend()
ax.set_ylim(0, 105)

# Add value labels
for bar in bars1 + bars2:
    height = bar.get_height()
    ax.annotate(f'{height:.1f}%', xy=(bar.get_x() + bar.get_width() / 2, height),
                xytext=(0, 3),
                textcoords="offset points",
                ha='center', va='bottom')

plt.tight_layout()
plt.savefig('rbf_vs_mlp_comparison.png', dpi=150)
plt.show()
```


=====

COMPARISON: RBF vs MLP on Iris Dataset

=====

Metric	RBF	MLP
Training Accuracy	96.67%	90.00%
Test Accuracy	93.33%	80.00%
Hidden Neurons	30	8
Training Method	Pseudoinverse	Backprop
Activation (Hidden)	Gaussian	Sigmoid
Activation (Output)	Linear	Sigmoid

RBF vs MLP Performance on Iris Dataset



Key Differences in Practice

Aspect	RBF	MLP
Training speed	Fast (two-stage, no iteration)	Slower (many epochs of gradient descent)
Interpretability	Higher (prototypes are actual data points)	Lower (weights are abstract)
Memory	Stores prototypes explicitly	Compressed into weights
Extrapolation	Poor (Gaussian falls to 0 far from prototypes)	Better (sigmoid doesn't decay to 0)
Best for	Local patterns, function approximation	Complex decision boundaries

6. Using sklearn (Library Alternative)

sklearn doesn't have a direct RBF network class, but provides related functionality:

1. **SVC with RBF kernel** - Support Vector Classifier using RBF kernel
2. **RBFSampler** - Approximates RBF kernel features

The SVC with RBF kernel is conceptually similar - it also uses Gaussian functions to measure similarity.

```
In [ ]: from sklearn.svm import SVC
from sklearn.kernel_approximation import RBFSampler
from sklearn.linear_model import SGDClassifier
from sklearn.pipeline import make_pipeline

# Method 1: SVC with RBF kernel
# This uses the RBF (Gaussian) kernel to measure similarity between samples

print("sklearn: SVC with RBF Kernel")
print("=" * 40)

svc_rbf = SVC(kernel='rbf', gamma='scale', random_state=42)
svc_rbf.fit(X_train, np.argmax(y_train, axis=1))

svc_train_accuracy = svc_rbf.score(X_train, np.argmax(y_train, axis=1))
svc_test_accuracy = svc_rbf.score(X_test, y_test_labels)

print(f"Number of RBF kernels: {len(svc_rbf.support_)}")
print(f"Training Accuracy: {svc_train_accuracy * 100:.2f}%")
print(f"Test Accuracy: {svc_test_accuracy * 100:.2f}%")
```

```
sklearn: SVC with RBF Kernel
=====
Number of RBF kernels: 38
Training Accuracy: 97.50%
Test Accuracy: 96.67%
```

```
In [ ]: # Method 2: RBFSampler + Linear Classifier
# This explicitly creates RBF features, similar to our from-scratch approach

print("\nsklearn: RBFSampler + Linear Classifier")
print("=" * 40)

# Create pipeline: RBF features → Linear classifier
rbf_pipeline = make_pipeline(
    RBFSampler(gamma=1.0, n_components=30, random_state=42), # 30 RBF features
    SGDClassifier(max_iter=1000, random_state=42)
)

rbf_pipeline.fit(X_train, np.argmax(y_train, axis=1))

rbf_pipe_train_accuracy = rbf_pipeline.score(X_train, np.argmax(y_train, axis=1))
rbf_pipe_test_accuracy = rbf_pipeline.score(X_test, y_test_labels)

print(f"Training Accuracy: {rbf_pipe_train_accuracy * 100:.2f}%")
print(f"Test Accuracy: {rbf_pipe_test_accuracy * 100:.2f}%")
```

```
sklearn: RBFSampler + Linear Classifier
=====
Training Accuracy: 95.83%
Test Accuracy: 96.67%
```

```
In [ ]: # Summary comparison

print("\n" + "=" * 60)
print("SUMMARY: All Methods Compared")
print("=" * 60)

methods = [
    ('Our RBF (from scratch)', test_accuracy),
    ('Our MLP (from scratch)', mlp_test_accuracy),
    ('sklearn SVC + RBF kernel', svc_test_accuracy),
    ('sklearn RBFSampler + SGD', rbf_pipe_test_accuracy),
]

print(f"\n{'Method':<30} {'Test Accuracy':>15}")
```

```

print("-" * 45)
for method, acc in methods:
    print(f"{method:<30} {acc*100:.1f}%")

print("\nNote: sklearn's SVC with RBF kernel is a powerful classifier that")
print("uses the same Gaussian similarity concept as RBF networks.")

```

```
=====
SUMMARY: All Methods Compared
=====
```

Method	Test Accuracy
Our RBF (from scratch)	93.33%
Our MLP (from scratch)	80.00%
sklearn SVC + RBF kernel	96.67%
sklearn RBFSampler + SGD	96.67%

Note: sklearn's SVC with RBF kernel is a powerful classifier that uses the same Gaussian similarity concept as RBF networks.

7. Experimentation Exercises

Now that you've seen how RBF networks work, try these experiments:

Exercise 7.1: Vary the Number of Prototypes

Try different numbers of prototypes (5, 10, 20, 50) and observe the effect on accuracy.

```
In [ ]: # Experiment: Effect of number of prototypes

prototype_counts = [5, 10, 15, 20, 30, 50]
results = []

print("Effect of Number of Prototypes")
print("=" * 40)

for n_proto in prototype_counts:
    rbf_exp = RBFNetwork(n_prototypes=n_proto)
    rbf_exp.fit(X_train, y_train)

    train_acc = accuracy_score(np.argmax(y_train, axis=1), rbf_exp.predict(X_train))
    test_acc = accuracy_score(y_test_labels, rbf_exp.predict(X_test))

    results.append((n_proto, train_acc, test_acc))
    print(f"Prototypes: {n_proto:3d} | Train: {train_acc*100:.1f}% | Test: {test_acc*100:.1f}%")

# Plot results
fig, ax = plt.subplots(figsize=(10, 5))

proto, train_accs, test_accs = zip(*results)
ax.plot(proto, [a*100 for a in train_accs], 'o-', label='Training', color='steelblue')
ax.plot(proto, [a*100 for a in test_accs], 's-', label='Test', color='coral')

ax.set_xlabel('Number of Prototypes')
ax.set_ylabel('Accuracy (%)')
ax.set_title('Effect of Number of Prototypes on RBF Performance')
ax.legend()
ax.grid(True, alpha=0.3)
ax.set_ylim(80, 102)

plt.tight_layout()
```

```

plt.savefig('rbf_prototype_experiment.png', dpi=150)
plt.show()

print("\nObservations:")
print("- Too few prototypes: underfitting (network can't represent the data)")
print("- Too many prototypes: overfitting risk and slower computation")
print("- Sweet spot depends on dataset complexity")

```

Effect of Number of Prototypes

Training RBF Network with 5 prototypes...

- Stage 1: Selected 5 prototypes via k-means
- Stage 2: Computed β values (mean=26.281)
- Stage 3: Learned output weights, shape=(6, 3)

Training complete!

Prototypes: 5 | Train: 90.8% | Test: 83.3%

Training RBF Network with 10 prototypes...

- Stage 1: Selected 10 prototypes via k-means
- Stage 2: Computed β values (mean=38.954)
- Stage 3: Learned output weights, shape=(11, 3)

Training complete!

Prototypes: 10 | Train: 95.8% | Test: 100.0%

Training RBF Network with 15 prototypes...

- Stage 1: Selected 15 prototypes via k-means
- Stage 2: Computed β values (mean=46.825)
- Stage 3: Learned output weights, shape=(16, 3)

Training complete!

Prototypes: 15 | Train: 97.5% | Test: 100.0%

Training RBF Network with 20 prototypes...

- Stage 1: Selected 20 prototypes via k-means
- Stage 2: Computed β values (mean=67.105)
- Stage 3: Learned output weights, shape=(21, 3)

Training complete!

Prototypes: 20 | Train: 96.7% | Test: 93.3%

Training RBF Network with 30 prototypes...

- Stage 1: Selected 30 prototypes via k-means
- Stage 2: Computed β values (mean=90.610)
- Stage 3: Learned output weights, shape=(31, 3)

Training complete!

Prototypes: 30 | Train: 96.7% | Test: 93.3%


Training RBF Network with 50 prototypes...

- Stage 1: Selected 50 prototypes via k-means
- Stage 2: Computed β values (mean=136.400)
- Stage 3: Learned output weights, shape=(51, 3)

Training complete!

Prototypes: 50 | Train: 98.3% | Test: 96.7%

Effect of Number of Prototypes on RBF Performance



Observations:

- Too few prototypes: underfitting (network can't represent the data)
- Too many prototypes: overfitting risk and slower computation
- Sweet spot depends on dataset complexity

Exercise 7.2: Experiment with β Values

Instead of computing β from the data, try using fixed β values and see the effect.

```
In [ ]: # Experiment: Effect of fixed β values

class RBFNetworkFixedBeta(RBFNetwork):
    """RBF Network with fixed (not computed) beta value."""

    def __init__(self, n_prototypes=10, fixed_beta=1.0):
        super().__init__(n_prototypes)
        self.fixed_beta = fixed_beta

    def fit(self, X, y):
        # Stage 1: Select prototypes (same as before)
        kmeans = KMeans(n_clusters=self.n_prototypes, random_state=42, n_init=10)
        kmeans.fit(X)
        self.prototypes = kmeans.cluster_centers_

        # Stage 2: Use fixed beta instead of computing from data
        self.betas = np.full(self.n_prototypes, self.fixed_beta)

        # Stage 3: Learn weights (same as before)
        H = self._compute_hidden_activations(X)
        H_bias = np.column_stack([H, np.ones(H.shape[0])])
        self.weights = np.linalg.pinv(H_bias) @ y

    return self

beta_values = [0.1, 0.5, 1.0, 2.0, 5.0, 10.0]
beta_results = []

print("Effect of Fixed β Values")
print("=" * 40)

for beta in beta_values:
    rbf_beta = RBFNetworkFixedBeta(n_prototypes=30, fixed_beta=beta)
    rbf_beta.fit(X_train, y_train)

    test_acc = accuracy_score(y_test_labels, rbf_beta.predict(X_test))
    beta_results.append((beta, test_acc))
    print(f"β = {beta:.5f} | Test Accuracy: {test_acc*100:.1f}%")

# Plot
fig, ax = plt.subplots(figsize=(10, 5))

betas, accs = zip(*beta_results)
ax.plot(betas, [a*100 for a in accs], 'o-', color='green', linewidth=2, markersize=8)

ax.set_xlabel('β Value')
ax.set_ylabel('Test Accuracy (%)')
ax.set_title('Effect of β on RBF Performance')
ax.set_xscale('log')
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig('rbf_beta_experiment.png', dpi=150)
plt.show()
```

```

print("\nObservations:")
print("- Small β: Wide Gaussians, neurons respond to distant inputs (underfitting)")
print("- Large β: Narrow Gaussians, neurons only respond to very close inputs (overfitting)")
print("- Computing β from data (as we did originally) usually gives good results")


```

Effect of Fixed β Values

```

β = 0.1 | Test Accuracy: 100.0%
β = 0.5 | Test Accuracy: 100.0%
β = 1.0 | Test Accuracy: 96.7%
β = 2.0 | Test Accuracy: 96.7%
β = 5.0 | Test Accuracy: 96.7%
β = 10.0 | Test Accuracy: 100.0%

```



Observations:

- Small β: Wide Gaussians, neurons respond to distant inputs (underfitting)
 - Large β: Narrow Gaussians, neurons only respond to very close inputs (overfitting)
 - Computing β from data (as we did originally) usually gives good results
-

Summary

In this self-study notebook, you have learned:

1. RBF Architecture: Three layers (input → RBF hidden → output) with Gaussian activation

2. Key Concepts:

- Prototypes: Training samples stored by RBF neurons
- Gaussian activation: $\varphi(x) = \exp(-\beta||x-p||^2)$
- β parameter: Controls width of Gaussian response
- Local response: Neurons only activate for nearby inputs

3. Training Process:

- Stage 1: Select prototypes (k-means clustering)
- Stage 2: Compute β values (from cluster spread)
- Stage 3: Learn output weights (pseudoinverse)

4. Comparison with MLP:

- RBF uses local Gaussian activation; MLP uses global sigmoid

- RBF training is faster (no iterative backpropagation)
- RBF is more interpretable (prototypes are actual data points)

5. **sklearn Alternatives:**

- SVC with RBF kernel
 - RBFSampler for explicit RBF features
-

Further Reading

- Lecture slides on RBF Networks
- Bishop, C.M. (1995). Neural Networks for Pattern Recognition. Oxford University Press.
- Haykin, S. (1999). Neural Networks: A Comprehensive Foundation. Prentice Hall.